

POCKETLEDGER

Your money, your model, your phone.

The private, on-device AI finance notebook the post-Mint generation is asking for.

Platform strategy & technical architecture

Built from the NotebookLM-for-finance signal | Q3 2026

THE SIGNAL

One XDA headline hid a market.

August 2025: a tech writer dragged bank statements into a research chatbot, called it brilliant, and the article crossed the personal-finance community in a weekend.

8 min

From install to first grounded answer about the user's own money

70%+

Of Gen Z say they are more careful with money than they used to be (Deloitte 2025)

84%

Of Gen Z and 79% of Millennials are familiar with generative AI (MEXC 2026)

THE WORKFLOW THAT WENT VIRAL

Eight minutes from panic to a model that knows your money.

1

Upload sources

Bank statements, credit-card PDFs, brokerage summaries -- all as documents.

2

Add goals

A second doc: "reduce discretionary spend, build a buffer, invest more."

3

Add frameworks

Articles on budgeting, investing, debt -- so the model knows your lens.

4

Ask in plain English

"Where am I overspending?" "What changed this month?"

A consumer research tool, in eight minutes, replaced her budgeting app. That is not a productivity hack. That is a category signal.

Three vectors converge in the same place.

01

The post-Mint vacuum

Intuit killed Mint on March 23, 2024. 3.6M users scattered. Monarch grew 20x, but every successor inherited a dashboard, not a dialogue. The user still has to interpret the chart.

02

Gen Z: mobile-first + AI-fluent + financially anxious

66% bank on mobile. 80%+ say money causes stress. 84% are familiar with gen AI. 70% are more careful than a year ago. They want a tool that explains, not one that displays.

03

The privacy + aggregation paradox

1 in 2 US adults have linked via Plaid. The 2022 \$58M class action is public. The same user who plugs Plaid into Monarch hand-redacts their bank statement before uploading it to NotebookLM.

Seven things the workflow does right.

- ✓ Upload-and-ask in eight minutes -- fast time-to-first-answer.
- ✓ Privacy is opt-in and explicit. No Plaid link by default.
- ✓ Grounded answers with citations. No hallucination about the user's data.
- ✓ Recurring check-in cadence feels natural -- once a month, refresh, ask.
- ✓ Plain English replaces dashboard navigation.
- ✓ Chat persists inside a bounded notebook. Bounded memory is the right memory.
- ✓ Goals and frameworks uploaded as separate sources.

None of this is in the existing finance apps. All of it is in the demand.

WHAT USERS HATE

Seven reasons the workflow does not survive contact with money.

01 The math is wrong

No Python environment. RAG-on-chunks hallucinates averages and sums.

02 It pretends to be mobile

Audio capped at 20 minutes, three per day. Two voices. No quick-log.

03 Wrong data substrate

No live feed, no OCR, no voice. 50 sources / 500K words per source.

04 Privacy is performative

Consumer tier trains on prompts. Workspace tier has telemetry. No on-device path.

05 No recurring ritual

No monthly check-in template. No subscription detector. No goal progress view.

06 Dense, non-shareable output

Text-heavy. No shareable report. Audio is chatty, not precise.

07 Black-box audit trail

You see the cited passage. You do not see chunk boundaries, retrieval scores, or rejection criteria.

THE GAP

The empty quadrant. Nobody owns it.

Privacy / No-bank-link

Plaid-linked

CONVERSATIONAL

POCKETLEDGER LIVES HERE

**Conversational + private + grounded
in your own papers**

*NotebookLM has the chat. The finance apps have the math.
Nobody ships both on a phone, with no bank link.*

NotebookLM

Chat OK No-bank OK
Math X Mobile X Privacy X

DASHBOARD

Finny / Vera

Mobile OK No-bank OK
Chat X Citations X Ritual X

Monarch / Cleo / Rocket

Mobile OK Math OK Plaid-linked
No-bank X On-device privacy X

01

THE PRODUCT

PocketLedger

A private, on-device AI finance notebook for your phone.
The answer to the question the XDA article was trying to ask.

POSITIONING

Built for Maya. Works for Daniel and Priya too.

ONE-LINE PITCH

**Your money, your
model, your
phone.**

CATEGORY WE OWN

Personal finance notebook. Not a budgeting app. Not an AI coach. A new thing.

Maya, 26

Primary

Designer, \$42-68K, 1-2 cards, BNPL habit, privacy-conscious, AI-fluent.

Daniel, 41

Pro arc

Married, two kids, freelancer, mixes personal + business, needs audit-ready exports.

Priya, 19

Gen-Z arc

College, first card, parents send money, wants streaks without shame.

Seven pillars. One product.

01

The Notebook

Sources, pages, threads, goals, check-ins. The durable object of the user's life.

02

Three Inputs

Voice, photo, chat. Multimodal capture that fits a phone.

03

Three Outputs

Chat with citations, check-in card, audio recap. Every claim cited.

04

The Math Layer

Deterministic engine next to the LLM. Numbers are reproducible, not sampled.

05

Goals + Frameworks

Curated library + free-form goal upload. The model reconciles goals against actuals.

06

Privacy Architecture

On-device by default. Cloud opt-in, E2E encrypted. We cannot read the data.

07 Sharing + Social -- Watermarked cards for stories, couples mode, accountant export, public notebooks.

THE SYMMETRY

Three inputs. Three outputs. One model.

INPUTS

How money gets in

Voice

"Spent fourteen fifty on lunch." On-device STT + intent parser. <2s.

Photo

Receipt, statement, contract. On-device OCR + LLM extraction. <3s for receipts.

Chat

Long-form conversation. "Where am I overspending?"
"What changed in March?"

OUTPUTS

How answers come back

Chat + citations

Every claim points to a source line. Tap to see the original. The trust wedge.

Check-in card

1-page structured artifact. Top movers, anomalies, goal progress, aura.

Audio recap

2-minute spoken summary. On-device TTS. Designed for the morning commute.

"Show me the math" is one tap on every number. The math engine is a separate process, not a feature of the LLM.

From install to first real win.

Day 0

Install

28MB, no signup. "No account needed. Your data lives on this phone." Pick a vibe.

Day 1

First capture

Snap a receipt. 3 seconds later: "\$4.75 at Blue Bottle, today 8:42am, Coffee. Looks right?"

Day 3

First voice entry

Walking home: "Spent eighteen on Thai food." Swipe yes on the lock screen.

Day 7

First check-in

Sunday 7pm push. Top movers, anomalies, goal progress, aura. One tap to share.

Day 30

First real win

\$140 less than last month. Three things drove it. Friend downloads PocketLedger.

The check-in ritual is the habit. Once a Sunday for 8 weeks, the switching cost is enormous.

MONETIZATION & METRICS

Cheaper than the category. Funnel over ARPU in year one.

FREE

\$0

forever, no ads

- 1 notebook, 5 pages
- 30 days history
- 30 logs / day
- On-device model
- 1 curated framework
- Check-in card, no share

PRO

\$4.99

per month or \$39.99/yr

- Unlimited everything
- All 30 frameworks
- Plaid link (opt-in)
- Couples mode
- Accountant export (PDF + CSV)
- E2E cloud sync

FAMILY

\$7.99

per month, up to 5 devices

- One shared notebook
- Personal notebooks
- Kids mode
- Allowance + visibility
- E2E cloud sync
- Priority on-device model

NORTH STAR: Weekly Check-ins Completed (WCC). Anti-metric: bank accounts linked -- we want this low.

02

THE ARCHITECTURE

Five layers.

Local-first, online-second, model-hot-swappable, structurally private. Designed so the company cannot read user data even if it wanted to.

FIVE LAYERS, TOP-DOWN

Stacked by dependency, not by importance.



The orchestrator never knows which model is on the other side. The platform never knows what the user asked. The company never sees the data.

THE HYBRID CONTRACT

On-device by default. Cloud for what the local model cannot do.

ON-DEVICE (default, forever)

- ✓ Voice-to-text
- ✓ Photo OCR (receipts, statements)
- ✓ Transaction categorization
- ✓ Sum / group / delta / trend / projection
- ✓ Subscription detection
- ✓ Goal reconciliation
- ✓ Check-in composition
- ✓ Chat with citations, short context

CLOUD (Pro tier, opt-in, E2E)

- Long-context chat (multi-year)
- "What if" scenario modeling
- Complex contract Q&A (v1.5)
- Tax-prep handoff (v2.0)
- Crypto onramp summaries (v2.0)
- *Zero retention contract with provider*
- *No training, audit log on device*

Rule: if the task can be done by reading the user's own notebook deterministically, it stays on-device. Always.

PRIVACY ARCHITECTURE

The architecture enforces the privacy. Not the policy page.

KEY HIERARCHY

Master key

256-bit. Generated on install. Never leaves OS keystore (Secure Enclave / StrongBox / WebCrypto).

Notebook key

Derived via HKDF. Rotated on notebook delete.

Source key

Derived from notebook. Allows per-source revocation.

Cloud envelope key

One-time, wraps master for cross-device sync. We never see it unwrapped.

E2E SYNC FLOW (Pro tier, opt-in)

Device A

generates cloudEnvelopeKey, encrypts master key locally, uploads ciphertext only

Apple / Google

sees only ciphertext in their app-data scope. Cannot decrypt. Cannot read.

Device B

reads ciphertext, asks Device A for envelope key over local network or BLE

Both devices

hold master key locally. Cloud stores only ciphertext. Forever.

We cannot read the notebook. Apple and Google cannot read the notebook. GDPR right-to-erasure is delete-the-app.

CLOSING

A new category is hiding inside one headline.

1 The demand is real

Post-Mint, AI-fluent, financially anxious, mobile-first. NotebookLM surfaced the signal because it was the only tool that already existed.

2 The gap is empty

Conversational + grounded + on-device + mobile + no-bank-link + the math. Nobody owns this quadrant.

3 The moat is structural

On-device model hot-swap. Deterministic math engine. E2E envelope. Source graph + citations. Check-in ritual. None of this is easy to copy.

Next step: ship the spine in Q1. iOS + Android v1.0. On-device chat + math + voice + photo. Free forever. The category will follow.

THE LLM STACK

Five tiers. Four run on the user's phone. One is opt-in cloud.

T5	Embedding (bge-small / Gemma embed) Vector search over the source graph. 33M params, 130MB. Runs on NPU.	\$0
T4	Cloud LLM (Claude Haiku 4.5 + Sonnet 4.6) Pro opt-in. E2E encrypted. Zero retention. BAA. Deep analysis only.	\$0.07/ user/mo
T3	Bundled fallback: Phi-3 Mini 3.8B / Gemma 3 4B For older iPhones / Android without AICore. 4-bit quantized, 2.3-2.5GB.	\$0
T2	OS-provided: Apple Foundation Models 3B / Gemini Nano 3.25B Free from Apple / Google. Ships in the OS. Zero bundle cost.	\$0
T1	FunctionGemma 270M (the traffic controller) Routes queries to tools. 50 tok/s on phone. 0.3s TTFT. Fine-tuned on our 50-tool API.	\$0

HOW EASY IS IT TO USE?

Four key moments. Measured in seconds.

Day 0	Install to first insight 5s of app time, 85s of human. All on-device, all private.	90s
Day 7	First weekly check-in Top movers, anomalies, goal progress, aura, watermarked share.	12s
Day 14	First chat query "Where am I overspending?" Cited answer, tappable sources.	<1s to first token
Day 30	Accountant export PDF + CSV with receipt thumbnails. Faster than opening	8s for 30 days
Eight rules: zero taps to value, no forms, no dashboard before data, every number ships with its calculation, every claim is cited.		

THE COST MODEL

\$1.2M to ship. \$141K/mo at 100K MAU.

CAPEX (one-time, v1.0 ship)

Product design (2 designers x 6mo)	\$180K
iOS engineering (2 x 6mo)	\$300K
Android engineering (2 x 6mo)	\$300K
Backend (1 x 4mo)	\$100K
ML / on-device (1 x 6mo)	\$150K
QA (1 x 6mo)	\$90K
Design system + assets	\$40K
Legal (BAA, privacy, ToS)	\$25K

Total CapEx	\$1.2M
--------------------	---------------

OPEX (monthly, at 100K MAU)

Salaries (8 people, loaded)	\$80K
Marketing (UA + content)	\$50K
Cloud LLM (Pro opt-in)	\$7K
Plaid (Pro opt-in)	\$1K
Cloud infra (Dokploy)	\$0.3K
Tools (Sentry, ASO, Notion)	\$2K
Anthropic BAA (amortized)	\$1K
Apple/Google dev fees	\$0.03K

Total OpEx	\$141K
-------------------	---------------

Variable cost on the free tier is \$0. The on-device model runs on the user's phone. Cloud LLM is only for Pro deep analysis.

UNIT ECONOMICS

We beat the category on every metric except ARPU.

5.5x

LTV / CAC (Y3)

Industry healthy: 3-5x. We hit 5.5x.

7.3mo

CAC payback (Y3)

Industry healthy: <12mo. We hit 7.3.

94%

Gross margin (Y3)

On-device inference = near-zero variable cost.

8%

Free -> Pro conversion (Y3)

Monarch 5-7%, Rocket 11% (we win mid-funnel).

\$4.50

Blended CPI (US)

iOS Finance \$11.62 global, \$4.10 US Apple Search Ads.

0


Free-tier variable cost

Every chat on the free tier is free to us.

Comparison (Y3): Monarch \$9.50 ARPU / 80% margin / 8-12mo payback. YNAB \$9.08 / 85% / 6-9mo. PocketLedger \$5.10 / 94% / 7.3mo. We win on margin + payback; they win on ARPU.

THREE-YEAR P&L

Break-even in Year 4 at ~\$25M ARR. \$58.5M total capital required.



THE NUMBERS

	Y1	Y2	Y3	Y4
MAU (avg)	125K	600K	1.75M	3.5M
Pro subs	5K	36K	140K	420K
Net revenue	\$144K	\$1.5M	\$5.5M	\$25M
Gross margin	65%	50%	89%	96%
Net income	(\$1.5M)	(\$3.6M)	(\$3.6M)	+\$10M
LTV/CAC (Pro)	0.6x	2.4x	5.5x	10x
Payback (mo)	42	17	7.3	4.7

Path: Pre-seed \$500K -> Seed \$3M -> Series A \$15M -> Series B \$40M -> profitable at \$25M ARR. ~\$58.5M total raised.

THE STORY

On-device AI. Zero variable cost. 18-month window.

"A private, on-device AI finance notebook for the post-Mint, AI-native generation. \$58.5M to break-even at \$25M ARR. 8.5x MOIC at the median exit."

5.5x

LTV / CAC by Y3

Industry healthy 3-5x. We sit at the top.

94%

Gross margin by Y3

On-device inference = near-zero variable cost.

\$58.5M

Total capital to break-even

Four rounds. Four years. Profitable at \$25M ARR.

Next step: close the \$500K pre-seed, ship the iOS prototype in 8 weeks, capture the XDA signal by Q1 2027.